



# SOFTWARE AUDIT REPORT

for

## AVA LABS



Prepared By: Shuxiao Wang

Hangzhou, China

Jun. 10, 2020

## Document Properties

Client	AVA Labs
Title	Software Audit Report
Target	Avalanche Blockchain
Version	0.1
Author	Jeff Liu
Auditors	Edward Lo, Ruiyi Zhang, Xudong Shao
Reviewed by	Chiachih Wu
Approved by	Xuxian Jiang
Classification	Confidential

## Version Info

Version	Date	Author	Description
0.1	Jun. 10, 2020	Jeff Liu	Initial Draft

## Contact

For more information about this document and its contents, please contact PeckShield Inc.

Name	Shuxiao Wang
Phone	+86 173 6454 5338
Email	contact@peckshield.com

## Contents

<b>1</b>	<b>Introduction</b>	<b>5</b>
1.1	About Avalanche Blockchain . . . . .	5
1.2	About PeckShield . . . . .	6
1.3	Methodology . . . . .	6
1.3.1	Risk Model . . . . .	6
1.3.2	Fuzzing . . . . .	7
1.3.3	White-box Audit . . . . .	8
1.4	Disclaimer . . . . .	10
<b>2</b>	<b>Findings</b>	<b>12</b>
2.1	Finding Summary . . . . .	12
2.2	Key Findings . . . . .	13
<b>3</b>	<b>Detailed Results</b>	<b>15</b>
3.1	DoS Vulnerability in the VM Module - #1 . . . . .	15
3.2	DoS Vulnerability in the VM Module - #2 . . . . .	16
3.3	DoS Vulnerability in the VM Module - #3 . . . . .	18
3.4	DoS Vulnerability in the VM Module - #4 . . . . .	19
3.5	Out-of-Memory Issue in the VM Module . . . . .	21
3.6	Missing Sanity Check in the Avalanche Consensus Process - #1 . . . . .	23
3.7	Missing Sanity Check in the Avalanche Consensus Process - #2 . . . . .	24
3.8	DoS Vulnerability in the Avalanche Consensus Module - #1 . . . . .	29
3.9	DoS Vulnerability in the Avalanche Consensus Module - #2 . . . . .	32
3.10	DoS Vulnerability in the Avalanche Consensus Module - #3 . . . . .	35
3.11	DoS Vulnerability in the Avalanche Consensus Module - #4 . . . . .	37
3.12	Missing Sanity Check in the Bootstrapping Process . . . . .	41
3.13	Missing Sanity Check in the Snowman Consensus Process - #1 . . . . .	43
3.14	Missing Sanity Check in the Snowman Consensus Process - #2 . . . . .	45
3.15	Missing Sanity Check in the RPC Module . . . . .	48
3.16	DoS Vulnerability in the RPC Module . . . . .	49

---

3.17 DoS Vulnerability in the P2P Module . . . . .	50
3.18 DoS Vulnerability in the VM Module - #5 . . . . .	52
<b>4 Conclusion</b>	<b>57</b>
<b>References</b>	<b>58</b>



# 1 | Introduction

Given the opportunity to review the **Avalanche Blockchain** design document and related source code, we in this report outline our systematic method to evaluate potential security issues in the Avalanche Blockchain implementation, expose possible semantic inconsistencies between the source code and the design specification, and provide additional suggestions and recommendations for improvement. Our results show that the given branch of Avalanche Blockchain can be further improved due to the presence of several issues related to either security or performance. This document describes our audit results in detail.

## 1.1 About Avalanche Blockchain

Avalanche Blockchain (AVA) [1] is a public blockchain system designed by AVA Labs [2]. AVA is a high-throughput, finance-focused, open-source blockchain system supporting highly decentralized applications and financial primitives, also interoperable blockchains. It can achieve 4000+ TPS with less than 3 seconds transaction finality time, mainly because of its DAG-optimized consensus protocol, and utilizing interoperable subnets to facilitate high scalability.

The basic information of Avalanche Blockchain is shown in Table 1.1, and its Git repository and the commit hash value (of the audited branch) are in Table 1.2.

Table 1.1: Basic Information of Avalanche Blockchain

Item	Description
Issuer	AVA Labs
Website	<a href="https://avalabs.org">https://avalabs.org</a>
Type	AVA Blockchain
Platform	Go
Audit Method	White-box
Latest Audit Report	Jun. 10, 2020

Table 1.2: The Commit Hash List Of Audited Branches

Git Repository	Commit Hash Of Audited Branch
<a href="https://github.com/ava-labs/gecko-security-analysis">github.com/ava-labs/gecko-security-analysis</a>	f66fd6de584e3e0347d39a62d427c440936ab745

## 1.2 About PeckShield

PeckShield Inc. [3] is a leading blockchain security company with the goal of elevating the security, privacy, and usability of current blockchain ecosystems by offering top-notch, industry-leading services and products including security audits. We are reachable at Telegram (<https://t.me/peckshield>), Twitter (<http://twitter.com/peckshield>), or Email ([contact@peckshield.com](mailto:contact@peckshield.com)).

## 1.3 Methodology

In the first phase of auditing Avalanche Blockchain, we use fuzzing to find out the corner cases that may not be covered by in-house testing. Next we do white-box auditing, in which PeckShield security auditors manually review Avalanche Blockchain design and source code, analyze them for any potential issues, and follow up with issues found in the fuzzing phase. If necessary, we design and implement individual test cases to further reproduce and verify the issues. In the following subsections, we will introduce the risk model as well as the audit procedure adopted in this report.

Table 1.3: Vulnerability Severity Classification

Impact	High	Critical	High	Medium
	Medium	High	Medium	Low
	Low	Medium	Low	Low
		High	Medium	Low
		Likelihood		

### 1.3.1 Risk Model

To standardize the evaluation, we define the following terminology based on OWASP Risk Rating Methodology [4]:

- Likelihood represents how likely a particular vulnerability is to be uncovered and exploited in the wild;
- Impact measures the technical loss and business damage of a successful attack;
- Severity demonstrates the overall criticality of the risk.

Likelihood and impact are categorized into three ratings: *H*, *M* and *L*, i.e., *high*, *medium* and *low* respectively. Severity is determined by likelihood and impact and can be classified into four categories accordingly, i.e., *Critical*, *High*, *Medium*, and *Low* shown in Table 1.3.

### 1.3.2 Fuzzing

Fuzzing or fuzz testing is an automated software testing technique of discovering software vulnerabilities by systematically finding and providing possible inputs to the target program, and then monitoring the program execution for crashes (or any unexpected results). In the first phase of our audit, we use fuzzing to find out possible corner cases or unusual inter-module interactions that may not be covered by in-house testing. As one of the most effective methods for exposing the presence of possible vulnerabilities, fuzzing technology has been the first choice for many security researchers in recent years. At present, there are many fuzzy testing tools and supporting software, which can help security personnels to conduct fuzzing and find vulnerabilities more efficiently. Based on the characteristics of the Avalanche Blockchain, we use AFL [5] as the primary tool for fuzz testing.

AFL (American Fuzzy Lop) is a security-oriented fuzzer that employs a novel type of compile-time instrumentation and genetic algorithms to automatically discover clean, interesting test cases that trigger new internal states in the targeted binary. Since its inception, AFL has gained growing popularity in the industry and has proved its effectiveness in discovering quite a few significant software bugs in a wide range of major software projects. The basic process of AFL fuzzing is as follows:

- Generate compile-time instrumentation to record information such as code execution path;
- Construct some input files to join the input queue, and change input files according to different strategies;
- Files that trigger a crash or timeout when executing an input file are logged for subsequent analysis;
- Loop through the above process.

Throughout the AFL testing, we will reproduce each crash based on the crash file generated by AFL. For each reported crash case, we will further analyze the root cause and check whether it is

indeed a vulnerability. Once a crash case is confirmed as a vulnerability of the Avalanche Blockchain, we will further analyze it as part of the white-box audit.

### 1.3.3 White-box Audit

After fuzzing, we continue the white-box audit by manually analyzing source code. Here we test target software's internal structure, design, coding, and we focus on verifying the flow of input and output through the application as well as examining possible design and implementation trade-offs for strengthened security. PeckShield auditors first fully review and understand the source code, then create specific test cases, execute them and analyze the results. Issues such as internal security loopholes, unexpected output, broken or poorly structured paths, etc., will be inspected under close scrutiny.

Blockchain is a secure method of creating a distributed database of transactions, and three major technologies of blockchain are cryptography, decentralization, and consensus model. Blockchain does come with unique security challenges, and based on our understanding of blockchain general design, we in this audit divide the blockchain software into the following major areas and inspect each area accordingly:

- Data and state storage, which is related to the database and files where blockchain data are saved.
- P2P networking, consensus, and transaction model in the networking layer. Note that the consensus and transaction logic is tightly coupled with networking.
- VM, account model, and incentive model. This is essentially the execution and business layer of the blockchain, and many blockchain business specific logics are implemented here.
- System contracts and services. These are system-level, blockchain-wide operation management contracts and services.
- Others. This includes any software modules that do not belong to above-mentioned areas, such as common crypto or other 3rd-party libraries, best practice or optimization used in other software projects, design and coding consistency, etc.

Based on the above classification, we show in Table 1.4 and Table 1.5 the detailed list of the audited items in this report.

To better describe each issue we identified, we also categorize the findings based on Common Weakness Enumeration (CWE-699) [6], which is a community-developed list of software weakness types to better classify and organize weaknesses around concepts frequently encountered in software development. We use the CWE categories in Table 1.6 to classify our findings.



Table 1.4: The Full List of Audited Items (Part I)

Category	Check Item
Data and State Storage	Blockchain Database Security
	Database State Integrity Check
Node Operation	Default Configuration Security
	Default Configuration Optimization
	Node Upgrade And Rollback Mechanism
Node Communication	External RPC Implementation Logic
	External RPC Function Security
	Node P2P Protocol Implementation Logic
	Node P2P Protocol Security
	Serialization/Deserialization
	Invalid/Malicious Node Management Mechanism
	Communication Encryption/Decryption
	Eclipse Attack Protection
	Fingerprint Attack Protection
Consensus	Consensus Algorithm Scalability
	Consensus Algorithm Implementation Logic
	Consensus Algorithm Security
Transaction Model	Transaction Privacy Security
	Transaction Fee Mechanism Security
	Transaction Congestion Attack Protection
VM	VM Implementation Logic
	VM Implementation Security
	VM Sandbox Escape
	VM Stack/Heap Overflow
	Contract Privilege Control
	Predefined Function Security
Account Model	Status Storage Algorithm Adjustability
	Status Storage Algorithm Security
	Double Spending Protection
Incentive Model	Mining Algorithm Security
	Mining Algorithm ASIC Resistance
	Tokenization Reward Mechanism

Table 1.5: The Full List of Audited Items (Part II)

Category	Check Item
System Contracts And Services	Memory Leak Detection
	Use-After-Free
	Null Pointer Dereference
	Undefined Behaviors
	Deprecated API Usage
	Signature Algorithm Security
	Multisignature Algorithm Security
SDK Security	Using RPC Functions Security
	Privatekey Algorithm Security
	Communication Security
	Function integrity checking code
Others	Third Party Library Security
	Memory Leak Detection
	Exception Handling
	Log Security
	Coding Suggestion And Optimization
	White Paper And Code Implementation Uniformity

## 1.4 Disclaimer

Note that this audit does not give any warranties on finding all possible security issues of the given blockchain software, i.e., the evaluation result does not guarantee the nonexistence of any further findings of security issues. As one audit cannot be considered comprehensive, we always recommend proceeding with several independent audits and a public bug bounty program to ensure the security of blockchain software. Last but not least, this security audit should not be used as an investment advice.

Table 1.6: Common Weakness Enumeration (CWE) Classifications Used in This Audit

Category	Summary
<b>Configuration</b>	Weaknesses in this category are typically introduced during the configuration of the software.
<b>Data Processing Issues</b>	Weaknesses in this category are typically found in functionality that processes data.
<b>Numeric Errors</b>	Weaknesses in this category are related to improper calculation or conversion of numbers.
<b>Security Features</b>	Weaknesses in this category are concerned with topics like authentication, access control, confidentiality, cryptography, and privilege management. (Software security is not security software)
<b>Time and State</b>	Weaknesses in this category are related to the improper management of time and state in an environment that supports simultaneous or near-simultaneous computation by multiple systems, processes, or threads.
<b>Error Conditions, Return Values, Status Codes</b>	Weaknesses in this category include weaknesses that occur if a function does not generate the correct return/status code, or if the application does not handle all possible return/status codes that could be generated by a function.
<b>Resource Management</b>	Weaknesses in this category are related to improper management of system resources.
<b>Behavioral Issues</b>	Weaknesses in this category are related to unexpected behaviors from code that an application uses.
<b>Business Logic</b>	Weaknesses in this category identify some of the underlying problems that commonly allow attackers to manipulate the business logic of an application. Errors in business logic can be devastating to an entire application.
<b>Initialization and Cleanup</b>	Weaknesses in this category occur in behaviors that are used for initialization and breakdown.
<b>Arguments and Parameters</b>	Weaknesses in this category are related to improper use of arguments or parameters within function calls.
<b>Expression Issues</b>	Weaknesses in this category are related to incorrectly written expressions within code.
<b>Input Validation Issues</b>	Weaknesses in this category are related to a software system's input validation components. Frequently these deal with sanitizing, neutralizing and validating any externally provided inputs to minimize malformed data from entering the system and preventing code injection in the input data.
<b>Coding Practices</b>	Weaknesses in this category are related to coding practices that are deemed unsafe and increase the chances that an exploitable vulnerability will be present in the application. They may not directly introduce a vulnerability, but indicate the product has not been carefully developed or maintained.




## 2 | Findings

### 2.1 Finding Summary

Here is a summary of our findings after analyzing Avalanche Blockchain. As mentioned earlier, we in the first phase of our audit studied AVA source code and ran our in-house static code analyzer through the codebase, and the purpose here is to statically identify known coding bugs, and then manually verify (reject or confirm) issues reported by our tools. After that, we manually review business logics, examine system operations, and place operation-specific aspects under scrutiny to uncover possible pitfalls and/or bugs.

We have so far identified a list of potential issues: some of them involve subtle corner cases that might not be previously thought of, while others refer to unusual interactions among multiple modules. For each uncovered issue, we have therefore developed test cases for reasoning, reproduction, and/or verification. After further analysis and internal discussion, we determined more than a dozen of issues of varying severities that need to be brought up, which are categorized in Table 2.1. More information can be found in the next subsection, and the detailed discussions of each of them are in Section 3.

Table 2.1: The Severity of Our Findings

Severity	# of Findings	
Critical	14	
High	2	
Medium	2	
Low	0	
Informational	0	
Total	18	

## 2.2 Key Findings

---

After analyzing all of the potential issues found during the audit, we determined that a number of them need to be brought up and paid more attention to, as shown in Table 2.2. Please refer to Section 3 for detailed discussion of each issue.



Table 2.2: Key Audit Findings

ID	Severity	Title	Category	Status
PVE-001	Critical	DoS Vulnerability in the VM Module - #1	Coding Practices	Fixed
PVE-002	Critical	DoS Vulnerability in the VM Module - #2	Business Logic Errors	Confirmed
PVE-003	Critical	DoS Vulnerability in the VM Module - #3	Business Logic Errors	Confirmed
PVE-004	Critical	DoS Vulnerability in the VM Module - #4	Business Logic Errors	Confirmed
PVE-005	High	Out-of-Memory Issue in the VM Module	Coding Practices	Fixed
PVE-006	Critical	Missing Sanity Check in the Avalanche Consensus Process - #1	Coding Practices	Confirmed
PVE-007	Critical	Missing Sanity Check in the Avalanche Consensus Process - #2	Coding Practices	Confirmed
PVE-008	Critical	DoS Vulnerability in the Avalanche Consensus Module - #1	Coding Practices	Confirmed
PVE-009	Critical	DoS Vulnerability in the Avalanche Consensus Module - #2	Coding Practices	Confirmed
PVE-010	Critical	DoS Vulnerability in the Avalanche Consensus Module - #3	Behavioral Problems	Confirmed
PVE-011	Critical	DoS Vulnerability in the Avalanche Consensus Module - #4	Coding Practices	Confirmed
PVE-012	Critical	Missing Sanity Check in the Bootstrapping Process	Input Validation Issues	Fixed
PVE-013	Critical	Missing Sanity Check in the Snowman Consensus Process - #1	Coding Practices	Confirmed
PVE-014	Critical	Missing Sanity Check in the Snowman Consensus Process - #2	Input Validation Issues	Fixed
PVE-015	High	Missing Sanity Check in the RPC Module	Input Validation Issues	Confirmed
PVE-016	Medium	DoS Vulnerability in the RPC Module	Coding Practices	Confirmed
PVE-017	Critical	DoS Vulnerability in the P2P Module	Coding Practices	Confirmed
PVE-018	Medium	DoS Vulnerability in the VM Module - #5	Coding Practices	Confirmed

## 3 | Detailed Results

### 3.1 DoS Vulnerability in the VM Module - #1

- ID: PVE-001
- Severity: Critical
- Likelihood: High
- Impact: High
- Target: snow/validators/sampler.go
- Category: Coding Practices [7]
- CWE subcategory: CWE-20 [8]

#### Description

There is a vulnerability in the platformvm module of AVA, which could be exploited by attackers to perform DoS attack.

```

65 // SampleValidatorsArgs are the arguments for calling SampleValidators
66 type SampleValidatorsArgs struct {
67     Size int `json:"size"`
68 }
69
70 // SampleValidatorsReply are the results from calling Sample
71 type SampleValidatorsReply struct {
72     Validators []string `json:"validators"`
73 }
74
75 // SampleValidators returns a sampling of the list of current validators
76 func (service *Service) SampleValidators(_ *http.Request, args *SampleValidatorsArgs
77     , reply *SampleValidatorsReply) error {
78     service.vm.ctx.Log.Debug("Sample called with {Size = %d}", args.Size)
79     validatorSet := service.vm.Validators.Sample(args.Size)

```

Listing 3.1: vms/platformvm/service.go

This API allows clients to interact with the P-Chain (Platform Chain), which maintains AVA's validator set and handles blockchain creation.

Platform.SampleValidators responds with sample size validators from the validator set of the default subnet.

```

139 // Sample implements the Group interface.
140 func (s *Sampler) Sample(size int) [] Validator {
141     s.lock.Lock()
142     defer s.lock.Unlock()
143
144     return s.sample(size)
145 }
146
147 func (s *Sampler) sample(size int) [] Validator {
148     list := make([] Validator, size)[:0]

```

Listing 3.2: snow/validators/sampler.go

`sample` creates a slice for `SampleValidators`(line 148). However, it would lead to a panic (runtime error: makeslice: len out of range) when parameter `size` is negative.

**Recommendation** Add a sanity check for the parameter `size`.

## 3.2 DoS Vulnerability in the VM Module - #2

- ID: PVE-002
- Severity: Critical
- Likelihood: High
- Impact: High
- Target: vms/spchainvm/block.go
- Category: Business Logic Errors [9]
- CWE subcategory: CWE-770 [10]

### Description

There is a vulnerability in gecko's `spchainvm` module, which could be exploited by attackers to perform DoS attack against the AVA network.

```

155 // ParseBlock implements the snowman.ChainVM interface
156 func (vm *VM) ParseBlock(b [] byte) (snowman.Block, error) {
157     c := Codec{}
158     rawBlock, err := c.UnmarshalBlock(b)
159     if err != nil {
160         return nil, err
161     }
162     block := &LiveBlock{
163         vm:    vm,
164         block: rawBlock,
165     }
166     if err := block.VerifyBlock(); err != nil {
167         return nil, err
168     }
169     return block, vm.state.SetBlock(vm.baseDB, rawBlock.ID(), rawBlock)
170 }

```

Listing 3.3: snow/engine/snowman/transitive.go



ParseBlock calls `unmarshalBlock` to deserialize the block, then check the validity of it by `block.VerifyBlock()`.

```

148 // VerifyBlock the validity of this block
149 func (lb *LiveBlock) VerifyBlock() error {
150     if lb.verifiedBlock {
151         return lb.validity
152     }
153
154     lb.verifiedBlock = true
155     lb.validity = lb.block.verify(lb.vm.ctx, &lb.vm.factory)
156     return lb.validity
157 }

```

Listing 3.4: `vms/spchainvm/live_block.go`

```

40 func (b *Block) verify(ctx *snow.Context, factory *crypto.FactorySECP256K1R) error {
41     switch {
42     case b == nil:
43         return errInvalidNil
44     case b.id.IsZero():
45         return errInvalidID
46     case b.parentID.IsZero():
47         return errInvalidID
48     }
49
50     for _, tx := range b.txs {
51         if err := tx.verify(ctx, factory); err != nil {
52             return err
53         }
54     }
55     return nil
56 }

```

Listing 3.5: `vms/spchainvm/block.go`

However, there is no check on sequence of transactions in blocks. An attacker could craft lots of blocks with valid contents (same transactions but in different order). `UnmarshalBlock` would successfully deserialize the blocks and calculate a different id. In the end, `SetBlock` would store them into database, which could cause the victim nodes Out-of-Storage.

Furthermore, the victims would pass the blocks to other sampled validators by calling `PushQuery`, and the new victims would repeat this process, resulting in a DDoS attack eventually.

**Recommendation** Preventing nodes from flooding the network with blocks/vertices.

### 3.3 DoS Vulnerability in the VM Module - #3

- ID: PVE-003
- Severity: Critical
- Likelihood: High
- Impact: High
- Target: vms/platformvm/prefixed\_state.go
- Category: Business Logic Errors [9]
- CWE subcategory: CWE-770 [10]

#### Description

There is a vulnerability in gecko's platformvm module, which could be exploited by attackers to perform DoS attack against the AVA network.

AVA network defines the core communication format between AVA nodes. It uses the primitive serialization format for payload packing and Salticidae's message format.

A node will send a Put message in response to receiving a Get message for a container the node has access to.

```

67 func (t *Transitive) Put(vdr ids.ShortID, requestID uint32, blkID ids.ID, blkBytes
    []byte) {
68     blk, err := t.Config.VM.ParseBlock(blkBytes)
69     if err != nil {
70         t.Config.Context.Log.Warn("ParseBlock failed due to %s for block:\n%s",
71             err,
72             formatting.DumpBytes{Bytes: blkBytes})
73         t.GetFailed(vdr, requestID, blkID)
74         return
75     }
76
77     t.insertFrom(vdr, blk)
78 }

```

Listing 3.6: snow/engine/snowman/transitive.go

Put handles an incoming Put request from other validators, it will call VM.ParseBlock to decode the Container.

```

275 func (vm *VM) ParseBlock(b []byte) (snowman.Block, error) {
276     blk, err := vm.codec.UnmarshalBlock(b)
277     if err != nil {
278         return nil, err
279     }
280     if blk, err := vm.getBlock(blk.ID()); err == nil {
281         return blk, nil
282     }
283     vm.state.SetBlock(vm.baseDB, blk)
284     return blk, nil
285 }
286

```

```

287 // GetBlock implements the snowman.ChainVM interface
288 func (vm *VM) GetBlock(blkID ids.ID) (snowman.Block, error) { return vm.getBlock(blkID)
    }
289
290 func (vm *VM) getBlock(blkID ids.ID) (Block, error) {
291     if blk, exists := vm.currentBlocks[blkID.Key()]; exists {
292         return blk, nil
293     }
294     if blk := vm.state.Block(vm.baseDB, blkID); blk != nil {
295         return blk, nil
296     }
297     return nil, errMissingBlock
298 }

```

Listing 3.7: vms/platformvm/vm.go

```

48 func (s *prefixedState) SetBlock(db database.Database, block Block) {
49     s.state.SetBlock(db, s.uniqueID(block.ID(), blockID, s.block), block)
50 }

```

Listing 3.8: vms/platformvm/prefixed\_state.go

ParseBlock calls getBlock to look up stored blocks by blkID (line 291, 294), then save the returned block into database (line 283)

However, there is not enough check on incoming blocks. An attacker could craft lots of blocks with valid contents, UnmarshalBlock would successfully deserialize the blocks and SetBlock would store them into database, which could cause the victim nodes Out-of-Storage.

Furthermore, the victims would pass the blocks to other sampled validators by calling PushQuery, and the new victims would repeat this process, resulting in a DDoS attack eventually.

**Recommendation** Preventing nodes from flooding the network with blocks/vertices.

## 3.4 DoS Vulnerability in the VM Module - #4

- ID: PVE-004
- Severity: Critical
- Likelihood: High
- Impact: High
- Target: vms/spchainvm/block.go
- Category: Business Logic Errors [9]
- CWE subcategory: CWE-770 [10]

### Description

There is a vulnerability in gecko's spchainvm module, which could be exploited by attackers to perform DoS attack against the AVA network.

AVA network defines the core communication format between AVA nodes. It uses the primitive serialization format for payload packing and Salticidae's message format.

A node will send a Put message in response to receiving a Get message for a container the node has access to.

```

67     func (t *Transitive) Put(vdr ids.ShortID, requestID uint32, blkID ids.ID, blkBytes
        []byte) {
68         blk, err := t.Config.VM.ParseBlock(blkBytes)
69         if err != nil {
70             t.Config.Context.Log.Warn("ParseBlock failed due to %s for block:\n%s",
71                 err,
72                 formatting.DumpBytes{Bytes: blkBytes})
73             t.GetFailed(vdr, requestID, blkID)
74             return
75         }
76
77         t.insertFrom(vdr, blk)
78     }

```

Listing 3.9: snow/engine/snowman/transitive.go

Put handles an incoming Put request from other validators, it would call VM.ParseBlock to decode the Container.

```

155     // ParseBlock implements the snowman.ChainVM interface
156     func (vm *VM) ParseBlock(b []byte) (snowman.Block, error) {
157         c := Codec{}
158         rawBlock, err := c.UnmarshalBlock(b)
159         if err != nil {
160             return nil, err
161         }
162         block := &LiveBlock{
163             vm:      vm,
164             block:   rawBlock,
165         }
166         if err := block.VerifyBlock(); err != nil {
167             return nil, err
168         }
169         return block, vm.state.SetBlock(vm.baseDB, rawBlock.ID(), rawBlock)
170     }

```

Listing 3.10: snow/engine/snowman/transitive.go

```

40     func (b *Block) verify(ctx *snow.Context, factory *crypto.FactorySECP256K1R) error {
41         switch {
42             case b == nil:
43                 return errInvalidNil
44             case b.id.IsZero():
45                 return errInvalidID
46             case b.parentID.IsZero():
47                 return errInvalidID
48         }

```

```

49
50     for _, tx := range b.txs {
51         if err := tx.verify(ctx, factory); err != nil {
52             return err
53         }
54     }
55     return nil
56 }

```

Listing 3.11: vms/spchainvm/block.go

ParseBlock calls UnmarshalBlock to deserialize the block, then check the validity of it by VerifyBlock (line 166), and finally calls SetBlock to store it into database.

The problem is that there are not enough sanity checks on the incoming blocks. An attacker could craft lots of blocks with valid contents, UnmarshalBlock would successfully deserialize the blocks and VerifyBlock could be easily bypassed if a crafted block has a non-zero parentID and zero txs. In the end, SetBlock would store them into database, which could cause the victim nodes Out-of-Storage.

Furthermore, the victims would pass the blocks to other sampled validators by calling PushQuery, and the new victims would repeat this process, resulting in a DDoS attack eventually.

**Recommendation** Preventing nodes from flooding the network with blocks/vertices.

### 3.5 Out-of-Memory Issue in the VM Module

- ID: PVE-005
- Severity: High
- Likelihood: Medium
- Impact: High
- Target: vms/platformvm/staker.go
- Category: Coding Practices [7]
- CWE subcategory: CWE-20 [8]

#### Description

This is a vulnerability in gecko's platformvm module.

```

57 // Verify that this transaction is well formed
58 func (tx *AddStakerTx) Verify(networkID uint32) error {
59     switch {
60     case tx == nil:
61         return errNilTx
62     case tx.key != nil:
63         return nil // Only verify the transaction once
64     case tx.id.IsZero():
65         return errInvalidID
66     case tx.networkID != networkID:
67         return errWrongNetworkID
68     }

```

```

69
70     if err := tx.staker.Verify(); err != nil {
71         return err
72     }

```

Listing 3.12: vms/platformvm/add\_staker\_tx.go

AddStakerTx is a proposal to add staker to the staking.

```

60 // Verify that this staker is well formed
61 func (staker *Staker) Verify() error {
62     switch {
63     case staker == nil:
64         return errNilInvalid
65     case staker.id.IsZero():
66         return errInvalidID
67     case staker.destination.IsZero():
68         return errInvalidDestination
69     case staker.amount == 0:
70         return errStakeAmount
71     case staker.endTime <= staker.startTime:
72         return errStakeDuration
73     default:
74         return nil
75     }
76 }

```

Listing 3.13: vms/platformvm/staker.go

staker.Verify() checks whether staker.endTime is greater than staker.startTime (line 71) but ignores the restrictions on duration. Also, we should ensure the proposed staker amount is greater than a Minimum value.

Specifically, a malicious attacker could flood these validators by intentionally proposing AddStakerTx with a large staker.endTime and 1 amount. After such a transaction is accepted, the new staker would not be removed in the future.

Also, a malicious attacker could flood these validators by intentionally proposing AddStakerTx with a large staker.endTime and 1 amount. The validators would keep storing these txs and eventually cause them Out-of-Storage.

**Recommendation** Add sanity checks for these parameters.

## 3.6 Missing Sanity Check in the Avalanche Consensus Process - #1

- ID: PVE-006
- Severity: Critical
- Likelihood: High
- Impact: High
- Target: snow/engine/avalanche/polls.go
- Category: Coding Practices [7]
- CWE subcategory: CWE-20 [8]

### Description

There is a vulnerability in gecko's Avalanche consensus module, which could be exploited by attackers to compromise the AVA network consensus.

AVA network defines the core communication format between AVA nodes. It uses the primitive serialization format for payload packing and Salticidae's message format.

A node would send a Chits message in response to receiving a PullQuery or PushQuery message for a container the node has added to consensus.

```

156 // Chits implements the Engine interface
157 func (t *Transitive) Chits(vdr ids.ShortID, requestID uint32, votes ids.Set) {
158     if !t.bootstrapped {
159         t.Config.Context.Log.Warn("Dropping Chits due to bootstrapping")
160         return
161     }
162
163     v := &voter{
164         t:      t,
165         vdr:    vdr,
166         requestID: requestID,
167         response: votes,
168     }
169     voteList := votes.List()
170     for _, vote := range voteList {
171         if !t.reinsertFrom(vdr, vote) {
172             v.deps.Add(vote)
173         }
174     }
175
176     t.vtxBlocked.Register(v)
177 }

```

Listing 3.14: snow/engine/avalanche/transitive.go

When the vote is registered, `update` would be called and executes `polls.Vote` to get the final voting results.

```

51 func (p *polls) Vote(requestID uint32, vdr ids.ShortID, votes []ids.ID) (ids.
    UniqueBag, bool) {

```

```

52     p.log.All("Vote. requestId: %d. validatorID: %s.", requestId, vdr)
53     poll, exists := p.m[requestID]
54     p.log.All("Poll: %v", poll)
55     if !exists {
56         return nil, false
57     }
58
59     poll.Vote(votes)
60     if poll.Finished() {
61         p.log.All("Poll is finished")
62         delete(p.m, requestId)
63         p.numPolls.Set(float64(len(p.m))) // Tracks performance statistics
64         return poll.votes, true
65     }
66     p.m[requestID] = poll
67     return nil, false
68 }

```

Listing 3.15: snow/engine/avalanche/polls.go

However, there is no sanity check on whether the validator has already voted. The `vdr` argument which is supposed to be used for checking the duplicate votes is never used here.

And every time `poll.Vote(votes)` gets called, `p.numPending` would decrease by one. The poll would be finished if `p.numPending` reaches 0.

Specifically, an attacker could vote many times to finish the poll before the victim gets votes from other validators. Since `requestId` is decoded from the message directly, attackers could target any `requestId` it wants to interrupt.

**Recommendation** Add a sanity check to forbid duplicate votes.

## 3.7 Missing Sanity Check in the Avalanche Consensus Process - #2

- |                      |   |
|----------------------|---|
| • ID: PVE-007        | • Target: snow/consensus/avalanche/topological.go |
| • Severity: Critical | • Category: Coding Practices [7]                  |
| • Likelihood: High   | • CWE subcategory: CWE-20 [8]                     |
| • Impact: High       |   |

### Description

This is a vulnerability in gecko's Avalanche consensus module of gecko, which could be exploited by attackers to compromise the AVA network consensus.



AVA network defines the core communication format between AVA nodes. It uses the primitive serialization format for payload packing and Salticidae's message format.

A node would send a Put message in response to receiving a Get message for a container the node has access to.

```

80 func (t *Transitive) Put(vdr ids.ShortID, requestID uint32, vtxID ids.ID, vtxBytes
    []byte) {
81     t.Config.Context.Log.All("Put called for vertexID %s", vtxID)
82
83     if !t.bootstrapped {
84         t.bootstrapper.Put(vdr, requestID, vtxID, vtxBytes)
85         return
86     }
87
88     vtx, err := t.Config.State.ParseVertex(vtxBytes)
89     if err != nil {
90         t.Config.Context.Log.Warn("ParseVertex failed due to %s for block:\n%s",
91             err,
92             formatting.DumpBytes{Bytes: vtxBytes})
93         t.GetFailed(vdr, requestID, vtxID)
94         return
95     }
96     t.insertFrom(vdr, vtx)
97 }

```

Listing 3.16: snow/engine/avalanche/transitive.go

Put handles an incoming Put request from other validators, it would call State.ParseVertex to decode the Container

The ParseVertex function here simply checks whether the vertex parentIDs and transactions are sorted and unique. If so, the function insertFrom and insert would be called to pass the vertex to consensus engine.

ParseVertex simply checks whether the vertex parentIDs and transactions are sorted and unique. If so, insertFrom and insert would be called to pass the vertex to consensus engine. When the vertex has Unknown parents, node would try sendRequest(vdr, parent.ID()) to get it back.

```

211 func (t *Transitive) insertFrom(vdr ids.ShortID, vtx avalanche.Vertex) bool {
212     issued := true
213     vts := []avalanche.Vertex{vtx}
214     for len(vts) > 0 {
215         vtx := vts[0]
216         vts = vts[1:]
217
218         if t.Consensus.VertexIssued(vtx) {
219             continue
220         }
221         if t.pending.Contains(vtx.ID()) {
222             issued = false
223             continue

```

```

224     }
225
226     for _, parent := range vtx.Parents() {
227         if !parent.Status().Fetched() {
228             t.sendRequest(vdr, parent.ID())
229             issued = false
230         } else {
231             vts = append(vts, parent)
232         }
233     }
234
235     t.insert(vtx)
236 }
237 return issued
238 }

```

Listing 3.17: snow/engine/avalanche/transitive.go

```

264     for _, tx := range txs {
265         for _, dep := range tx.Dependencies() {
266             depID := dep.ID()
267             if !txIDs.Contains(depID) && !t.Consensus.TxIssued(dep) {
268                 t.missingTxns.Add(depID)
269                 i.txDeps.Add(depID)
270             }
271         }
272     }
273
274     t.Config.Context.Log.All("Vertex: %s is blocking on %d vertices and %d transactions",
275                             , vtxID, i.vtxDeps.Len(), i.txDeps.Len())
276
277     t.vtxBlocked.Register(&vtxIssuer{i: i})
278     t.txBlocked.Register(&txIssuer{i: i})

```

Listing 3.18: snow/engine/avalanche/transitive.go

As we can see here, two new issuers would be registered in `vtxBlocked` and `txBlocked`. If they have dependencies which haven't been issued, `t.vtxBlocked` would record that issuer to its corresponding `depID` key in the map. Finally, `update` would be called.

```

48     func (b *Blocker) Register(pending Blockable) {
49         b.init()
50
51         for _, pendingID := range pending.Dependencies().List() {
52             key := pendingID.Key()
53             (*b)[key] = append((*b)[key], pending)
54         }
55
56         pending.Update()
57     }

```

Listing 3.19: snow/events/blocker.go

```

35 func (i *issuer) Update() {
36     if i.abandoned || i.issued || i.vtxDeps.Len() != 0 || i.txDeps.Len() != 0 || i.t
        .Consensus.VertexIssued(i.vtx) {
37         return
38     }
39     i.issued = true
40
41     vtxID := i.vtx.ID()
42     i.t.pending.Remove(vtxID)
43
44     for _, tx := range i.vtx.Txs() {
45         if err := tx.Verify(); err != nil {
46             i.t.Config.Context.Log.Debug("Transaction failed verification due to %s,
                dropping vertex", err)
47             i.t.vtxBlocked.Abandon(vtxID)
48             return
49         }
50     }
51
52     i.t.Config.Context.Log.All("Adding vertex to consensus:\n%s", i.vtx)
53
54     i.t.Consensus.Add(i.vtx)
55
56     p := i.t.Consensus.Parameters()
57     vdrs := i.t.Config.Validators.Sample(p.K) // Validators to sample
58
59     vdrSet := ids.ShortSet{} // Validators to sample repr. as a set
60     for _, vdr := range vdrs {
61         vdrSet.Add(vdr.ID())
62     }
63
64     i.t.RequestID++
65     if numVdrs := len(vdrs); numVdrs == p.K && i.t.polls.Add(i.t.RequestID, vdrSet.
        Len()) {
66         i.t.Config.Sender.PushQuery(vdrSet, i.t.RequestID, vtxID, i.vtx.Bytes())
67     } else if numVdrs < p.K {
68         i.t.Config.Context.Log.Error("Query for %s was dropped due to an
            insufficient number of validators", vtxID)
69     }
70
71     i.t.vtxBlocked.Fulfill(vtxID)
72     for _, tx := range i.vtx.Txs() {
73         i.t.txBlocked.Fulfill(tx.ID())
74     }
75 }

```

Listing 3.20: snow/engine/avalanche/issuer.go

If `i.abandoned` or `i.vtxDeps.Len() != 0`, it would return immediately. When dependencies are resolved, node would call `i.t.vtxBlocked.Fulfill(vtxID)` to remove the corresponding ID in `i.vtxDeps`. In the end, the vertex would be added into consensus (line 54).

However, this `issuer` could be abandoned on purpose by intentionally triggering `GetFailed`.

```

100 func (t *Transitive) GetFailed(vdr ids.ShortID, requestID uint32, vtxID ids.ID) {
101     if !t.bootstrapped {
102         t.bootstrapper.GetFailed(vdr, requestID, vtxID)
103         return
104     }
105
106     t.pending.Remove(vtxID)
107     t.vtxBlocked.Abandon(vtxID)
108     t.vtxReqs.Remove(vtxID)
109
110     if t.vtxReqs.Len() == 0 {
111         for _, txID := range t.missingTxs.List() {
112             t.txBlocked.Abandon(txID)
113         }
114         t.missingTxs.Clear()
115     }
116
117     // Track performance statistics
118     t.numVtxRequests.Set(float64(t.vtxReqs.Len()))
119     t.numTxRequests.Set(float64(t.missingTxs.Len()))
120     t.numBlockedVtx.Set(float64(t.pending.Len()))
121 }

```

Listing 3.21: snow/engine/avalanche/transitive.go

Specifically, an attacker could send `Put` message with an invalid vertex with a target `vtxID`, and the following process would abandon the vertex. Since the `vtxID` is passed in by user directly, attackers could target any vertex he wants to abandon.

```

35 func (b *Blocker) Abandon(id ids.ID) {
36     b.init()
37
38     key := id.Key()
39     blocking := (*b)[key]
40     delete(*b, key)
41
42     for _, pending := range blocking {
43         pending.Abandon(id)
44     }
45 }

```

Listing 3.22: snow/consensus/avalanche/topological.go

Corresponding `convincer` would also be affected due to the chain reaction of abandon process. In the end, the node can not respond to a `PushQuery` message sent by other validators.

**Recommendation** Enhance the sanity check on the input vertices.

## 3.8 DoS Vulnerability in the Avalanche Consensus Module - #1

- ID: PVE-008
- Severity: Critical
- Likelihood: High
- Impact: High
- Target: snow/consensus/avalanche/topological.go
- Category: Coding Practices [7]
- CWE subcategory: CWE-20 [8]

### Description

This is a vulnerability in gecko's Avalanche consensus module, which could be exploited by attackers to perform DoS attack against the AVA network.

AVA network defines the core communication format between AVA nodes. It uses the primitive serialization format for payload packing and Salticidae's message format.

A node would send a Put message in response to receiving a Get message for a container the node has access to.

```

80 func (t *Transitive) Put(vdr ids.ShortID, requestID uint32, vtxID ids.ID, vtxBytes
    []byte) {
81     t.Config.Context.Log.All("Put called for vertexID %s", vtxID)
82
83     if !t.bootstrapped {
84         t.bootstrapper.Put(vdr, requestID, vtxID, vtxBytes)
85         return
86     }
87
88     vtx, err := t.Config.State.ParseVertex(vtxBytes)
89     if err != nil {
90         t.Config.Context.Log.Warn("ParseVertex failed due to %s for block:\n%s",
91             err,
92             formatting.DumpBytes{Bytes: vtxBytes})
93         t.GetFailed(vdr, requestID, vtxID)
94         return
95     }
96     t.insertFrom(vdr, vtx)
97 }

```

Listing 3.23: snow/engine/avalanche/transitive.go

Put handles an incoming Put request from other validators, it would call State.ParseVertex to decode the Container

The ParseVertex function here simply checks whether the vertex parentIDs and transactions are sorted and unique. If so, the function insertFrom and insert would be called to pass the vertex to consensus engine.

```

264 for _, tx := range txs {
265     for _, dep := range tx.Dependencies() {

```

```

266         depID := dep.ID()
267         if !txIDs.Contains(depID) && !t.Consensus.TxIssued(dep) {
268             t.missingTxs.Add(depID)
269             i.txDeps.Add(depID)
270         }
271     }
272 }
273
274 t.Config.Context.Log.All("Vertex: %s is blocking on %d vertices and %d transactions"
    , vtxID, i.vtxDeps.Len(), i.txDeps.Len())
275
276 t.vtxBlocked.Register(&vtxIssuer{i: i})
277 t.txBlocked.Register(&txIssuer{i: i})

```

Listing 3.24: snow/engine/avalanche/transitive.go

As we can see here, two new issuers would be registered in `vtxBlocked` and `txBlocked`. When they are registered, their `Update` function would be called.

```

35 func (i *issuer) Update() {
36     if i.abandoned || i.issued || i.vtxDeps.Len() != 0 || i.txDeps.Len() != 0 || i.t
        .Consensus.VertexIssued(i.vtx) {
37         return
38     }
39     i.issued = true
40
41     vtxID := i.vtx.ID()
42     i.t.pending.Remove(vtxID)
43
44     for _, tx := range i.vtx.Txs() {
45         if err := tx.Verify(); err != nil {
46             i.t.Config.Context.Log.Debug("Transaction failed verification due to %s,
                dropping vertex", err)
47             i.t.vtxBlocked.Abandon(vtxID)
48             return
49         }
50     }
51
52     i.t.Config.Context.Log.All("Adding vertex to consensus:\n%s", i.vtx)
53
54     i.t.Consensus.Add(i.vtx)

```

Listing 3.25: snow/engine/avalanche/issuer.go

`i.t.Consensus.Add(i.vtx)` would be called after passing the semantic legitimacy of `vtx`'s transactions.

```

112 func (ta *Topological) Add(vtx Vertex) {
113     ta.ctx.Log.AssertTrue(vtx != nil, "Attempting to insert nil vertex")
114
115     key := vtx.ID().Key()
116     if vtx.Status().Decided() {
117         return // Already decided this vertex

```

```

118     } else if _, exists := ta.nodes[key]; exists {
119         return // Already inserted this vertex
120     }
121
122     for _, tx := range vtx.Txs() {
123         if !tx.Status().Decided() {
124             // Add the consumers to the conflict graph.
125             ta.cg.Add(tx)
126         }
127     }
128
129     ta.nodes[key] = vtx // Add this vertex to the set of nodes
130     ta.numProcessing.Inc()
131
132     ta.update(vtx) // Update the vertex and it's ancestry
133 }

```

Listing 3.26: snow/consensus/avalanche/topological.go

Add calls `ta.update` to accept the vertex. Finally, `db.Commit` would store it into database.

```

360 // Check my parent statuses
361 for _, dep := range deps {
362     if status := dep.Status(); status == choices.Rejected {
363         vtx.Reject() // My parent is rejected, so I should be rejected
364         ta.numRejected.Inc()
365         delete(ta.nodes, vtxKey)
366         ta.numProcessing.Dec()
367
368         ta.preferenceCache[vtxKey] = false
369         ta.virtuousCache[vtxKey] = false
370         return
371     } else if status != choices.Accepted {
372         acceptable = false // My parent isn't accepted, so I can't be
373     }
374 }

```

Listing 3.27: snow/consensus/avalanche/topological.go

However, there is not enough sanity check on the input vertices.

Specifically, an attacker could craft lots of vertices with valid contents, all the sanity checks could be easily bypassed if a crafted vertex has nil txs and different `Rejected` parents. In the end, `db.Commit` would store them into database(line 363), which could cause the victim nodes Out-of-Storage.

Furthermore, the victims would pass the vertices to other sampled validators by calling `PushQuery`, and the new victims would repeat this process, resulting in a DDoS attack eventually.

**Recommendation** Enhance the sanity check on the incoming vertices.

## 3.9 DoS Vulnerability in the Avalanche Consensus Module - #2

- ID: PVE-009
- Severity: Critical
- Likelihood: High
- Impact: High
- Target: snow/consensus/avalanche/topological.go
- Category: Coding Practices [7]
- CWE subcategory: CWE-20 [8]

### Description

There is a vulnerability in gecko's Avalanche consensus module, which could be exploited by attackers to perform DoS attack against the AVA network.

AVA network defines the core communication format between AVA nodes. It uses the primitive serialization format for payload packing and Salticidae's message format.

A node would send a Put message in response to receiving a Get message for a container the node has access to.

```

80 func (t *Transitive) Put(vdr ids.ShortID, requestID uint32, vtxID ids.ID, vtxBytes
    []byte) {
81     t.Config.Context.Log.All("Put called for vertexID %s", vtxID)
82
83     if !t.bootstrapped {
84         t.bootstrapper.Put(vdr, requestID, vtxID, vtxBytes)
85         return
86     }
87
88     vtx, err := t.Config.State.ParseVertex(vtxBytes)
89     if err != nil {
90         t.Config.Context.Log.Warn("ParseVertex failed due to %s for block:\n%s",
91             err,
92             formatting.DumpBytes{Bytes: vtxBytes})
93         t.GetFailed(vdr, requestID, vtxID)
94         return
95     }
96     t.insertFrom(vdr, vtx)
97 }

```

Listing 3.28: snow/engine/avalanche/transitive.go

Put handles an incoming Put request from other validators, it would call State.ParseVertex to decode the Container

The ParseVertex function here simply checks whether the vertex parentIDs and transactions are sorted and unique. If so, the function insertFrom and insert would be called to pass the vertex to consensus engine.

```

264 for _, tx := range txs {
265     for _, dep := range tx.Dependencies() {

```



```

266         depID := dep.ID()
267         if !txIDs.Contains(depID) && !t.Consensus.TxIssued(dep) {
268             t.missingTxs.Add(depID)
269             i.txDeps.Add(depID)
270         }
271     }
272 }
273
274 t.Config.Context.Log.All("Vertex: %s is blocking on %d vertices and %d transactions"
    , vtxID, i.vtxDeps.Len(), i.txDeps.Len())
275
276 t.vtxBlocked.Register(&vtxIssuer{i: i})
277 t.txBlocked.Register(&txIssuer{i: i})

```

Listing 3.29: snow/engine/avalanche/transitive.go

As we can see here, two new issuers would be registered in `vtxBlocked` and `txBlocked`. When they are registered, their `Update` function would be called.

```

35 func (i *issuer) Update() {
36     if i.abandoned || i.issued || i.vtxDeps.Len() != 0 || i.txDeps.Len() != 0 || i.t
        .Consensus.VertexIssued(i.vtx) {
37         return
38     }
39     i.issued = true
40
41     vtxID := i.vtx.ID()
42     i.t.pending.Remove(vtxID)
43
44     for _, tx := range i.vtx.Txs() {
45         if err := tx.Verify(); err != nil {
46             i.t.Config.Context.Log.Debug("Transaction failed verification due to %s,
                dropping vertex", err)
47             i.t.vtxBlocked.Abandon(vtxID)
48             return
49         }
50     }
51
52     i.t.Config.Context.Log.All("Adding vertex to consensus:\n%s", i.vtx)
53
54     i.t.Consensus.Add(i.vtx)

```

Listing 3.30: snow/engine/avalanche/issuer.go

`i.t.Consensus.Add(i.vtx)` will be called after passing the semantic legitimacy of `vtx`'s transactions.

```

112 func (ta *Topological) Add(vtx Vertex) {
113     ta.ctx.Log.AssertTrue(vtx != nil, "Attempting to insert nil vertex")
114
115     key := vtx.ID().Key()
116     if vtx.Status().Decided() {
117         return // Already decided this vertex
118     } else if _, exists := ta.nodes[key]; exists {

```

```

119         return // Already inserted this vertex
120     }
121
122     for _, tx := range vtx.Txs() {
123         if !tx.Status().Decided() {
124             // Add the consumers to the conflict graph.
125             ta.cg.Add(tx)
126         }
127     }
128
129     ta.nodes[key] = vtx // Add this vertex to the set of nodes
130     ta.numProcessing.Inc()
131
132     ta.update(vtx) // Update the vertex and it's ancestry
133 }

```

Listing 3.31: snow/consensus/avalanche/topological.go

Add calls `ta.update` to accept the vertex. Finally, `db.Commit` will store it into database.

```

412     switch {
413     case acceptable:
414         // I'm acceptable, why not accept?
415         vtx.Accept()
416         ta.numAccepted.Inc()
417         delete(ta.nodes, vtxKey)
418         ta.numProcessing.Dec()
419     case rejectable:
420         // I'm rejectable, why not reject?
421         vtx.Reject()
422         ta.numRejected.Inc()
423         delete(ta.nodes, vtxKey)
424         ta.numProcessing.Dec()
425     }

```

Listing 3.32: snow/consensus/avalanche/topological.go

Although there is a lack of sanity check on the vertices input, to be exact, an attacker could craft lots of vertices with valid contents, `parseVertex` would successfully deserialize them and `vtx.Verify` can be easily bypassed if:

1. a crafted vertex has nil `parentID` and nil `txs` but different `height` will be `Accept()` (line 415).
2. a crafted vertex contains different rejected `txs` will be `Reject()` (line 421).

In the end, `db.Commit` will store them into database, which could cause the victim nodes Out-of-Storage.

Furthermore, the victims would pass the vertices to other sampled validators by calling `PushQuery`, and the new victims would repeat this process, resulting in a DDoS attack eventually.

**Recommendation** Enhance the sanity check on the vertices input.

## 3.10 DoS Vulnerability in the Avalanche Consensus Module - #3

- ID: PVE-010
- Severity: Critical
- Likelihood: High
- Impact: High
- Target: snow/engine/avalanche/transitive.go
- Category: Bad Coding Practices [7]
- CWE subcategory: CWE-20 [8]

### Description

There is a vulnerability in gecko's Avalanche consensus module, which could be exploited by attackers to perform DoS attack against the AVA network.

AVA network defines the core communication format between AVA nodes. It uses the primitive serialization format for payload packing and Salticidae's message format.

A node would send a Put message in response to receiving a Get message for a container the node has access to.

```

80 func (t *Transitive) Put(vdr ids.ShortID, requestID uint32, vtxID ids.ID, vtxBytes []
    byte) {
81     t.Config.Context.Log.All("Put called for vertexID %s", vtxID)
82
83     if !t.bootstrapped {
84         t.bootstrapper.Put(vdr, requestID, vtxID, vtxBytes)
85         return
86     }
87
88     vtx, err := t.Config.State.ParseVertex(vtxBytes)
89     if err != nil {
90         t.Config.Context.Log.Warn("ParseVertex failed due to %s for block:\n%s",
91             err,
92             formatting.DumpBytes{Bytes: vtxBytes})
93         t.GetFailed(vdr, requestID, vtxID)
94         return
95     }
96     t.insertFrom(vdr, vtx)
97 }
```

Listing 3.33: snow/engine/avalanche/transitive.go

Put handles an incoming Put request from other validators, it will call State.ParseVertex to decode the Container.

```

60 func (s *Serializer) ParseVertex(b []byte) (avacon.Vertex, error) {
61     vtx, err := s.parseVertex(b)
62     if err != nil {
```

```

63     return nil, err
64 }
65 if err := vtx.Verify(); err != nil {
66     return nil, err
67 }
68 uVtx := &uniqueVertex{
69     serializer: s,
70     vtxID:      vtx.ID(),
71 }
72 if uVtx.Status() == choices.Unknown {
73     uVtx.setVertex(vtx)
74 }
75
76 s.db.Commit()
77 return uVtx, nil
78 }

```

Listing 3.34: snow/engine/avalanche/state/serializer.go

```

41 func (vtx *vertex) Verify() error {
42     switch {
43     case !ids.IsSortedAndUniqueIDs(vtx.parentIDs):
44         return errInvalidParents
45     case !ids.IsSortedAndUniqueTxns(vtx.txns):
46         return errInvalidTxns
47     default:
48         return nil
49     }
50 }

```

Listing 3.35: snow/engine/avalanche/state/vertex.go

ParseVertex calls parseVertex to deserialize the vertex, then check the validity of it by vtx.Verify, which will make sure parentIDs and txns are sorted and unique. Finally, db.Commit will store it into database.

Although there is a lack of sanity check on the vertices input, to be exact, an attacker could craft lots of vertices with valid contents, parseVertex would successfully deserialize them and vtx.Verify can be easily bypassed if a crafted vertex has unique and sorted parentID / txns (or none of these at all). In the end, db.Commit would store them into database, which could cause the victim nodes Out-of-Storage.

Furthermore, the victims would pass the vertices to other sampled validators by calling PushQuery, and the new victims would repeat this process, resulting in a DDoS attack eventually.

**Recommendation** Add check on vertices with accepted transactions.

### 3.11 DoS Vulnerability in the Avalanche Consensus Module - #4

- ID: PVE-011
- Severity: Critical
- Likelihood: High
- Impact: High
- Target: snow/engine/avalanche/transitive.go
- Category: Bad Coding Practices [7]
- CWE subcategory: CWE-20 [8]

#### Description

There is a vulnerability in gecko's Avalanche consensus module, which could be exploited by attackers to perform DoS attack against the AVA network.

AVA network defines the core communication format between AVA nodes. It uses the primitive serialization format for payload packing and Salticidae's message format.

A node would send a Put message in response to receiving a Get message for a container the node has access to.

```

80 func (t *Transitive) Put(vdr ids.ShortID, requestID uint32, vtxID ids.ID, vtxBytes []
    byte) {
81     t.Config.Context.Log.All("Put called for vertexID %s", vtxID)
82
83     if !t.bootstrapped {
84         t.bootstrapper.Put(vdr, requestID, vtxID, vtxBytes)
85         return
86     }
87
88     vtx, err := t.Config.State.ParseVertex(vtxBytes)
89     if err != nil {
90         t.Config.Context.Log.Warn("ParseVertex failed due to %s for block:\n%s",
91             err,
92             formatting.DumpBytes{Bytes: vtxBytes})
93         t.GetFailed(vdr, requestID, vtxID)
94         return
95     }
96     t.insertFrom(vdr, vtx)
97 }

```

Listing 3.36: snow/engine/avalanche/transitive.go

Put handles incoming Put messages, it would use State.ParseVertex to decode the Container. ParseVertex simply checks whether the vertex parentIDs and transactions are sorted and unique. If so, insertFrom and insert will be called to pass the vertex to consensus engine.

```

211 func (t *Transitive) insertFrom(vdr ids.ShortID, vtx avalanche.Vertex) bool {

```

```

212 issued := true
213 vts := []avalanche.Vertex{vtx}
214 for len(vts) > 0 {
215     vtx := vts[0]
216     vts = vts[1:]
217
218     if t.Consensus.VertexIssued(vtx) {
219         continue
220     }
221     if t.pending.Contains(vtx.ID()) {
222         issued = false
223         continue
224     }
225
226     for _, parent := range vtx.Parents() {
227         if !parent.Status().Fetched() {
228             t.sendRequest(vdr, parent.ID())
229             issued = false
230         } else {
231             vts = append(vts, parent)
232         }
233     }
234
235     t.insert(vtx)
236 }
237 return issued
238 }

```

Listing 3.37: snow/engine/avalanche/transitive.go

```

240 func (t *Transitive) insert(vtx avalanche.Vertex) {
241     vtxID := vtx.ID()
242
243     t.pending.Add(vtxID)
244     t.vtxReqs.Remove(vtxID)
245
246     i := &issuer{
247         t: t,
248         vtx: vtx,
249     }
250
251     for _, parent := range vtx.Parents() {
252         if !t.Consensus.VertexIssued(parent) {
253             i.vtxDeps.Add(parent.ID())
254         }
255     }
256
257     txs := vtx.Txs()
258
259     txIDs := ids.Set{}
260     for _, tx := range txs {
261         txIDs.Add(tx.ID())
262     }

```

```

263
264     for _, tx := range txs {
265         for _, dep := range tx.Dependencies() {
266             depID := dep.ID()
267             if !txIDs.Contains(depID) && !t.Consensus.TxIssued(dep) {
268                 t.missingTxns.Add(depID)
269                 i.txDeps.Add(depID)
270             }
271         }
272     }
273
274     t.Config.Context.Log.All("Vertex: %s is blocking on %d vertices and %d transactions",
275                             vtxID, i.vtxDeps.Len(), i.txDeps.Len())
276
277     t.vtxBlocked.Register(&vtxIssuer{i: i})
278     t.txBlocked.Register(&txIssuer{i: i})
279
280     if t.vtxReqs.Len() == 0 {
281         for _, txID := range t.missingTxns.List() {
282             t.txBlocked.Abandon(txID)
283         }
284         t.missingTxns.Clear()
285     }
286
287     // Track performance statistics
288     t.numVtxRequests.Set(float64(t.vtxReqs.Len()))
289     t.numTxRequests.Set(float64(t.missingTxns.Len()))
290     t.numBlockedVtx.Set(float64(t.pending.Len()))
291 }

```

Listing 3.38: snow/engine/avalanche/transitive.go

As we can see here, if the vertex's parents haven't been fetched, it would call `sendRequest` to get the parent vertex. And two new issuers would be registered in `vtxBlocked` and `txBlocked`. When registered, `Update` will be called.

```

35 func (i *issuer) Update() {
36     if i.abandoned || i.issued || i.vtxDeps.Len() != 0 || i.txDeps.Len() != 0 || i.t.
37         Consensus.VertexIssued(i.vtx) {
38         return
39     }
40     i.issued = true
41
42     vtxID := i.vtx.ID()
43     i.t.pending.Remove(vtxID)
44
45     for _, tx := range i.vtx.Txns() {
46         if err := tx.Verify(); err != nil {
47             i.t.Config.Context.Log.Debug("Transaction failed verification due to %s, dropping
48                 vertex", err)
49             i.t.vtxBlocked.Abandon(vtxID)
50         }
51     }
52     return
53 }

```

```

49     }
50 }
51
52 i.t.Config.Context.Log.All("Adding vertex to consensus:\n%s", i.vtx)
53
54 i.t.Consensus.Add(i.vtx)
55
56 p := i.t.Consensus.Parameters()
57 vdrs := i.t.Config.Validators.Sample(p.K) // Validators to sample
58
59 vdrSet := ids.ShortSet{} // Validators to sample repr. as a set
60 for _, vdr := range vdrs {
61     vdrSet.Add(vdr.ID())
62 }
63
64 i.t.RequestID++
65 if numVdrs := len(vdrs); numVdrs == p.K && i.t.polls.Add(i.t.RequestID, vdrSet.Len())
66 {
67     i.t.Config.Sender.PushQuery(vdrSet, i.t.RequestID, vtxID, i.vtx.Bytes())
68 } else if numVdrs < p.K {
69     i.t.Config.Context.Log.Error("Query for %s was dropped due to an insufficient number
70     of validators", vtxID)
71 }
72
73 i.t.vtxBlocked.Fulfill(vtxID)
74 for _, tx := range i.vtx.Txs() {
75     i.t.txBlocked.Fulfill(tx.ID())
76 }
77 }

```

Listing 3.39: snow/engine/avalanche/issuer.go

If sendRequest timed out, GetFailed will clear the registered issuers.

```

100 func (t *Transitive) GetFailed(vdr ids.ShortID, requestID uint32, vtxID ids.ID) {
101     if !t.bootstrapped {
102         t.bootstrapper.GetFailed(vdr, requestID, vtxID)
103         return
104     }
105
106     t.pending.Remove(vtxID)
107     t.vtxBlocked.Abandon(vtxID)
108     t.vtxReqs.Remove(vtxID)
109
110     if t.vtxReqs.Len() == 0 {
111         for _, txID := range t.missingTxs.List() {
112             t.txBlocked.Abandon(txID)
113         }
114         t.missingTxs.Clear()
115     }
116
117     // Track performance statistics
118     t.numVtxRequests.Set(float64(t.vtxReqs.Len()))

```



```

119 t.numTxRequests.Set(float64(t.missingTxs.Len()))
120 t.numBlockedVtx.Set(float64(t.pending.Len()))
121 }

```

Listing 3.40: snow/engine/avalanche/transitive.go

However, if the length of vtxReqs is not equal to zero, the corresponding issuers in txBlocked won't be deleted and missingTxs won't be cleared.

Specifically, an attacker could send lots of vertices with non-existent vertex parents to flood the txBlocked and missingTxs. Eventually, this would cause the node Out-of-Memory.

**Recommendation** Delete the corresponding txBlocked and missingTxs after timeout.

### 3.12 Missing Sanity Check in the Bootstrapping Process

- ID: PVE-012
- Severity: Critical
- Likelihood: High
- Impact: High
- Target: snow/engine/avalanche/bootstrapper.go
- Category: Input Validation Issues [11]
- CWE subcategory: CWE-349 [12]

#### Description

This is a vulnerability in the consensus module of gecko, which could be exploited by attackers to compromise the AVA network consensus.

AVA network defines the core communication format between AVA nodes. It uses the primitive serialization format for payload packing and Salticidae's message format.

A node would send a Put message in response to receiving a Get message for a container the node has access to.

```

87 func (b *bootstrapper) Put(vdr ids.ShortID, requestID uint32, vtxID ids.ID, vtxBytes []
    byte) {
88     b.BootstrapConfig.Context.Log.All("Put called for vertexID %s", vtxID)
89
90     if !b.pending.Contains(vtxID) {
91         return
92     }
93
94     vtx, err := b.State.ParseVertex(vtxBytes)
95     if err != nil {
96         b.BootstrapConfig.Context.Log.Warn("ParseVertex failed due to %s for block:\n%s",
97             err,
98             formatting.DumpBytes{Bytes: vtxBytes})
99         b.GetFailed(vdr, requestID, vtxID)

```

```

100     return
101 }
102
103 b.addVertex(vtx)
104 }

```

Listing 3.41: snow/engine/avalanche/bootstrapper.go

ParseVertex would try to extract the vertex from vtxBytes, and the vtx.status would be set to Processing if extracted successfully. After that, addVertex would create vertexJob / txJob for the received vertex and transactions, and push them to VtxBlocked and TxBlocked as long as the vtxID is in pending.

At last, executeAll would execute the tx / vertex job, to accept or reject them.

```

199 func (b *bootstrapper) executeAll(jobs *queue.Jobs, numBlocked prometheus.Gauge) {
200     for job, err := jobs.Pop(); err == nil; job, err = jobs.Pop() {
201         numBlocked.Dec()
202         if err := jobs.Execute(job); err != nil {
203             b.BootstrapConfig.Context.Log.Warn("Error executing: %s", err)
204         }
205     }
206 }

```

Listing 3.42: snow/engine/avalanche/bootstrapper.go

jobs.Execute execute the job's Execute, for example, txJob.Execute for the transactions.

```

44 func (t *txJob) Execute() {
45     if t.MissingDependencies().Len() != 0 {
46         t.numDropped.Inc()
47         return
48     }
49
50     switch t.tx.Status() {
51     case choices.Unknown, choices.Rejected:
52         t.numDropped.Inc()
53     case choices.Processing:
54         t.tx.Accept()
55         t.numAccepted.Inc()
56     }
57 }

```

Listing 3.43: snow/engine/avalanche/tx\_job.go

However, in the case of choices.Processing, transactions could be accepted without any sanity checks (line 54), e.g., tx.Verify for the validity of the transaction.

Specifically, an attacker could send Put messages with a vtxID in pending and an illegal vertex to other nodes, and the illegal transaction / vertex would be accepted unconditionally since there is no sanity checks at all.

**Recommendation** Add check whether blkID corresponds to blkBytes in Put function.

### 3.13 Missing Sanity Check in the Snowman Consensus Process

#### - #1

- ID: PVE-013
- Severity: Critical
- Likelihood: High
- Impact: High
- Target: snow/engine/snowman/polls.go
- Category: Coding Practices [7]
- CWE subcategory: CWE-20 [8]

#### Description

This is a vulnerability in gecko's Snowman consensus module, which could be exploited by attackers to compromise the AVA network consensus.

AVA network defines the core communication format between AVA nodes. It uses the primitive serialization format for payload packing and Salticidae's message format.

A node would send a Chits message in response to receiving a PullQuery or PushQuery message for a container the node has added to consensus.

```

112 // Chits implements the Engine interface
113 func (t *Transitive) Chits(vdr ids.ShortID, requestID uint32, votes ids.Set) {
114     // Since this is snowman, there should only be one ID in the vote set
115     if votes.Len() != 1 {
116         t.Config.Context.Log.Warn("Chits was called with the wrong number of votes %
117             d. RequestID: %d", votes.Len(), requestID)
118         t.QueryFailed(vdr, requestID)
119         return
120     }
121     vote := votes.List()[0]
122     t.Config.Context.Log.All("Chit was called. RequestID: %v. Vote: %s", requestID,
123         vote)
124     v := &voter{
125         t:      t,
126         vdr:    vdr,
127         requestID: requestID,
128         response: vote,
129     }
130
131     if !t.reinsertFrom(vdr, vote) {
132         v.deps.Add(vote)
133     }
134
135     t.blocked.Register(v)
136 }

```

Listing 3.44: snow/engine/snowman/transitive.go

```

24 func (v *voter) Update() {
25     if v.deps.Len() != 0 {
26         return
27     }
28
29     results := ids.Bag{}
30     finished := false
31     if v.response.IsZero() {
32         results, finished = v.t.polls.CancelVote(v.requestID, v.vdr)
33     } else {
34         results, finished = v.t.polls.Vote(v.requestID, v.vdr, v.response)
35     }
36
37     if !finished {
38         return
39     }
40
41     v.t.Config.Context.Log.All("Finishing poll [%d] with:\n%s", v.requestID, &
42         results)
43     v.t.Consensus.RecordPoll(results)
44
45     v.t.Config.VM.SetPreference(v.t.Consensus.Preference())
46
47     if v.t.Consensus.Finalized() {
48         v.t.Config.Context.Log.All("Snowman engine can quiesce")
49         return
50     }
51     v.t.Config.Context.Log.All("Snowman engine can't quiesce")
52
53     if len(v.t.polls.m) == 0 {
54         v.t.repoll()
55     }
56 }

```

Listing 3.45: snow/engine/snowman/voter.go

When the vote is registered, `Update` would be called and executes `polls.Vote` for voting (line 34), or cancel the vote (line 32).

```

36 func (p *polls) Vote(requestID uint32, vdr ids.ShortID, vote ids.ID) (ids.Bag, bool)
37 {
38     p.log.All("[polls.Vote] Vote: requestID: %d. validatorID: %s. Vote: %s",
39         requestID, vdr, vote)
40     poll, exists := p.m[requestID]
41     if !exists {
42         return ids.Bag{}, false
43     }
44     poll.Vote(vote)
45     if poll.Finished() {
46         delete(p.m, requestID)
47         p.numPolls.Set(float64(len(p.m))) // Tracks performance statistics
48         return poll.votes, true
49     }
50 }

```

```

47     }
48     p.m[requestID] = poll
49     return ids.Bag{}, false
50 }
51
52 // CancelVote registers the connections failure to respond to a query for [id].
53 func (p *polls) CancelVote(requestID uint32, vdr ids.ShortID) (ids.Bag, bool) {
54     p.log.All("CancelVote received. requestID: %d. validatorID: %s. Vote: %s",
55         requestID, vdr)
56     poll, exists := p.m[requestID]
57     if !exists {
58         return ids.Bag{}, false
59     }
60     poll.CancelVote()
61     if poll.Finished() {
62         delete(p.m, requestID)
63         p.numPolls.Set(float64(len(p.m))) // Tracks performance statistics
64         return poll.votes, true
65     }
66     p.m[requestID] = poll
67     return ids.Bag{}, false
68 }

```

Listing 3.46: snow/engine/snowman/polls.go

However, neither of these two has any sanity checks on whether the validator has already voted. The `vdr` argument which is supposed to be used for checking the duplicate votes is never used here.

Specifically, an attacker could vote many times to finish the poll before the victim gets votes from other validators. Since `requestID` is decoded from the message directly, attackers could target any `requestID` it wants to interrupt.

**Recommendation** Add a sanity check to forbid duplicate votes.

### 3.14 Missing Sanity Check in the Snowman Consensus Process

#### - #2

- ID: PVE-014
- Severity: Critical
- Likelihood: High
- Impact: High
- Target: snow/engine/snowman/transitive.go
- Category: Input Validation Issues [11]
- CWE subcategory: CWE-349 [12]

## Description

There is a vulnerability in gecko's Snowman consensus module, which could be exploited by attackers to compromise the AVA network consensus.

AVA network defines the core communication format between AVA nodes. It uses the primitive serialization format for payload packing and Salticidae's message format.

A node would send a Put message in response to receiving a Get message for a container the node has access to.

```

67 func (t *Transitive) Put(vdr ids.ShortID, requestID uint32, blkID ids.ID, blkBytes []
    byte) {
68     blk, err := t.Config.VM.ParseBlock(blkBytes)
69     if err != nil {
70         t.Config.Context.Log.Warn("ParseBlock failed due to %s for block:\n%s",
71             err,
72             formatting.DumpBytes{Bytes: blkBytes})
73         t.GetFailed(vdr, requestID, blkID)
74         return
75     }
76
77     t.insertFrom(vdr, blk)
78 }

```

Listing 3.47: snow/engine/snowman/transitive.go

Put handles incoming Put messages, it would call VM.ParseBlock to decode the Container.

insertFrom and insert would be called if the block can be unmarshalled. If the block's parent isn't fetched yet, the node would try sendRequest(vdr, parentID) to get it back.

```

174 func (t *Transitive) insertFrom(vdr ids.ShortID, blk snowman.Block) bool {
175     blkID := blk.ID()
176     for !t.Consensus.Issued(blk) && !t.pending.Contains(blkID) {
177         t.insert(blk)
178
179         parent := blk.Parent()
180         parentID := parent.ID()
181         if parentStatus := parent.Status(); !parentStatus.Fetched() {
182             t.sendRequest(vdr, parentID)
183             return false
184         }
185
186         blk = parent
187         blkID = parentID
188     }
189     return !t.pending.Contains(blkID)
190 }

```

Listing 3.48: snow/engine/snowman/transitive.go

```

200 func (t *Transitive) insert(blk snowman.Block) {

```

```

201 blkID := blk.ID()
202
203 t.pending.Add(blkID)
204 t.blkReqs.Remove(blkID)
205
206 i := &issuer{
207     t:    t,
208     blk: blk,
209 }
210
211 if parent := blk.Parent(); !t.Consensus.Issued(parent) {
212     parentID := parent.ID()
213     t.Config.Context.Log.All("Block waiting for parent %s", parentID)
214     i.deps.Add(parentID)
215 }
216
217 t.blocked.Register(i)
218
219 // Tracks performance statistics
220 t.numBlkRequests.Set(float64(t.blkReqs.Len()))
221 t.numBlockedBlk.Set(float64(t.pending.Len()))
222 }

```

Listing 3.49: snow/engine/snowman/transitive.go

As we can see, the new issuer would be registered to blocked. If they have dependencies which haven't been issued, t.blocked would record that issuer to its corresponding depID keys in the map. Finally, update would be called.

```

48 func (b *Blocker) Register(pending Blockable) {
49     b.init()
50
51     for _, pendingID := range pending.Dependencies().List() {
52         key := pendingID.Key()
53         (*b)[key] = append((*b)[key], pending)
54     }
55
56     pending.Update()
57 }

```

Listing 3.50: snow/events/blocker.go

```

35 func (i *issuer) Update() {
36     if i.abandoned || i.deps.Len() != 0 {
37         return
38     }
39
40     i.t.deliver(i.blk)
41 }

```

Listing 3.51: snow/engine/snowman/issuer.go

If `i.abandoned` or `i.deps.Len() != 0`, it would return immediately. When dependencies are resolved, `deliver` would be called to add the block into consensus and send it out (line 40).

However, this issuer could be abandoned on purpose by intentionally triggering `GetFailed`.

Specifically, an attacker could send `Put` message with an invalid block which will fail `VM.ParseBlock` and the following process would abandon the block. Since the block ID is passed in by user directly, attacker can target any block he wants to abandon.

```

35 func (b *Blocker) Abandon(id ids.ID) {
36     b.init()
37
38     key := id.Key()
39     blocking := (*b)[key]
40     delete(*b, key)
41
42     for _, pending := range blocking {
43         pending.Abandon(id)
44     }
45 }

```

Listing 3.52: `snow/events/blocker.go`

Corresponding `convincer` would also be affected due to the chain reaction of abandon process. In the end, the node can not response to a `PushQuery` message sent by other validators.

**Recommendation** Add check whether `blkID` corresponds to `blkBytes` in `Put` function.

### 3.15 Missing Sanity Check in the RPC Module

- ID: PVE-015
- Severity: High
- Likelihood: High
- Impact: Medium
- Target: `api/admin/service.go`
- Category: Input Validation Issues [11]
- CWE subcategory: CWE-349 [12]

#### Description

This is a vulnerability in the RPC module of `gecko`, which could be exploited by attackers to overwrite arbitrary files on the node server if the node is running with root privilege.

Ava node provides RPC service for users to interact with the node. The `admin.startCPUProfiler` and `admin.memoryProfile` APIs are used to dump the profile information into the specified file. The node handles the `rpc` call with the function `StartCPUProfiler`.

```

108 func (service *Admin) StartCPUProfiler(r *http.Request, args *StartCPUProfilerArgs,
109     reply *StartCPUProfilerReply) error {
    service.log.Debug("Admin: StartCPUProfiler called with %s", args.Filename)

```



```

110     reply.Success = true
111     return service.performance.StartCPUProfiler(args.Filename)
112 }

```

Listing 3.53: api/admin/service.go

It would get the filename from the arguments and write the profile information into that file.

```

20 func (p *Performance) StartCPUProfiler(filename string) error {
21     if p.cpuProfileFile != nil {
22         return errCPUProfilerRunning
23     }
24
25     file, err := os.Create(filename)
26     if err != nil {
27         return err
28     }
29     if err := pprof.StartCPUProfile(file); err != nil {
30         file.Close()
31         return err
32     }
33     runtime.SetMutexProfileFraction(1)
34
35     p.cpuProfileFile = file
36     return nil
37 }

```

Listing 3.54: api/admin/performance.go

However, there is no sanity checks on the filename passed in, nor any privilege requirements on the rpc call.

Specifically, the attacker could send the target file path and file name to the rpc call (like ".././etc/passwd") and the file would be overwritten by the node depends on the node process's privilege.

**Recommendation** Add check on input file path to prevent traversal attack.

### 3.16 DoS Vulnerability in the RPC Module

- ID: PVE-016
- Severity: Medium
- Likelihood: High
- Impact: Low
- Target: vms/avm/vm.go
- Category: Coding Practices [7]
- CWE subcategory: CWE-20 [8]

## Description

There is a potential performance issue in the the X-Chain API of gecko, which could be exploited by attackers to perform a DoS attack against an AVA node.

```

243 func (vm *VM) GetUTXOs(addr IDs.Set) ([]*UTXO, error) {
244     utxoIDs := IDs.Set{}
245     for _, addr := range IDs.List() {
246         utxos, _ := vm.state.Funds(addr)
247         utxoIDs.Add(utxos...)
248     }
249
250     utxos := []*UTXO{}
251     for _, utxoID := range utxoIDs.List() {
252         utxo, err := vm.state.UTXO(utxoID)
253         if err != nil {
254             return nil, err
255         }
256         utxos = append(utxos, utxo)
257     }
258     return utxos, nil
259 }

```

Listing 3.55: vms/avm/vm.go

GetUTXOs returns the utxos that at least one of the provided addresses is referenced in. However, these UTXOs are stored in a single list for each account, and there is no limitation on the length of the passed in addr. Technically, attackers could send requests to this API with a series of addresses which may own numerous utxos, to exhaust the resources of the victim node.

**Recommendation** Limit the length of the input addr.

## 3.17 DoS Vulnerability in the P2P Module

- ID: PVE-017
- Severity: Critical
- Likelihood: High
- Impact: High
- Target: networking/handshake\_handlers.go
- Category: Coding Practices [7]
- CWE subcategory: CWE-20 [8]

## Description

This is a vulnerability in the P2P module of gecko, which could be exploited by attackers to perform DoS attack against the AVA network.

AVA network defines the core communication format between AVA nodes. It uses the primitive serialization format for payload packing and Salticidae's message format.

```

124 net.RegConnHandler(salticidae.MsgNetworkConnCallback(C.checkPeerCertificate), nil)
125 peerNet.RegPeerHandler(salticidae.PeerNetworkPeerCallback(C.peerHandler), nil)
126 peerNet.RegUnknownPeerHandler(salticidae.PeerNetworkUnknownPeerCallback(C.
    unknownPeerHandler), nil)
127 net.RegHandler(Ping, salticidae.MsgNetworkMsgCallback(C.ping), nil)
128 net.RegHandler(Pong, salticidae.MsgNetworkMsgCallback(C.pong), nil)
129 net.RegHandler(GetVersion, salticidae.MsgNetworkMsgCallback(C.getVersion), nil)
130 net.RegHandler(Version, salticidae.MsgNetworkMsgCallback(C.version), nil)
131 net.RegHandler(GetPeerList, salticidae.MsgNetworkMsgCallback(C.getPeerList), nil)
132 net.RegHandler(PeerList, salticidae.MsgNetworkMsgCallback(C.peerList), nil)

```

Listing 3.56: networking/handshake\_handlers.go

During node setup process, Initialize registers many handlers for different kind of messages.

```

502 func peerList(_msg *C.struct_msg_t, _conn *C.struct_msgnetwork_conn_t, _unsafe.
    Pointer) {
503     HandshakeNet.numPeerlistReceived.Inc()
504
505     msg := salticidae.MsgFromC(salticidae.CMsg(_msg))
506     build := Builder{}
507     pMsg, err := build.Parse(PeerList, msg.GetPayloadByMove())
508     if err != nil {
509         HandshakeNet.log.Warn("Failed to parse PeerList message due to %s", err)
510         // TODO: What should we do here?
511         return
512     }
513
514     ips := pMsg.Get(Peers).([]utils.IPDesc)
515     cErr := salticidae.NewError()
516     for _, ip := range ips {
517         HandshakeNet.log.All("Trying to adding peer %s", ip)
518         addr := salticidae.NewNetAddrFromIPPortString(ip.String(), false, &cErr)
519         if cErr.GetCode() == 0 && !HandshakeNet.myAddr.IsEq(addr) { // Make sure not
            to connect to myself
520             ip := toIPDesc(addr)
521
522             if !HandshakeNet.pending.ContainsIP(addr) && !HandshakeNet.connections.
                ContainsIP(addr) {
523                 HandshakeNet.log.Debug("Adding peer %s", ip)
524                 HandshakeNet.net.AddPeer(addr)
525             }
526         }
527         addr.Free()
528     }
529 }

```

Listing 3.57: networking/handshake\_handlers.go

peerlist is registered to handle peerList message. If the addr is not in pending or connections, it will call AddPeer to store the new peer.

```

497 func (self PeerNetwork) AddPeer(peer PeerID) int32 {

```

```

498     return int32(C.peernetwork_add_peer(self.inner, peer.inner))
499 }

```

Listing 3.58: network.go

The new peer would be added to `known_peers` which is a `std::unordered_map`.

However, there is no limitation nor any sanity checks on the passed in peers in the `peerlist` message.

Specifically, the attacker could flood the target with `peerlist` messages containing different peer addresses which could cause the victim nodes Out-of-Memory.

The same problem exists in `unknownPeerHandler`.

```

352 func unknownPeerHandler(_addr *C.netaddr_t, _cert *C.x509_t, _unsafe.Pointer) {
353     addr := salticidae.NetAddrFromC(salticidae.CNetAddr(_addr))
354     ip := toIPDesc(addr)
355     HandshakeNet.log.Info("Adding peer %s", ip)
356     HandshakeNet.net.AddPeer(addr)
357 }

```

Listing 3.59: networking/handshake\_handlers.go

In C++ library, if the `ping_handler` receives a ping message with unknown claimed address, it would call `unknown_peer_cb`.

`unknown_peer_cb` would call `unknownPeerHandler` to add the unknown peer to `HandshakeNet.net`. Since the `addr` is decoded from the ping message directly(`msg.claimed_addr`), the attacker could use a different addresses everytime.

Specifically, the attacker could flood the target with C++ level ping messages(`MsgPing`) containing different claimed addresses which could cause the victim nodes Out-of-Memory.

**Recommendation** Use LRU cache or add a length limitation on `known_peers`.

## 3.18 DoS Vulnerability in the VM Module - #5

- ID: PVE-018
- Severity: Medium
- Likelihood: Low
- Impact: High
- Target: `coreth/consensus/dummy/consensus.go`
- Category: Coding Practices [7]
- CWE subcategory: CWE-20 [8]

### Description

This is a vulnerability in the EVM module of gecko, which could be exploited by attackers to perform DoS attack against the AVA network.

AVA network defines the core communication format between AVA nodes. It uses the primitive serialization format for payload packing and Salticidae's message format.

A node would send a Put message in response to receiving a Get message for a container the node has access to.

```

67     func (t *Transitive) Put(vdr ids.ShortID, requestID uint32, blkID ids.ID, blkBytes
        []byte) {
68         blk, err := t.Config.VM.ParseBlock(blkBytes)
69         if err != nil {
70             t.Config.Context.Log.Warn("ParseBlock failed due to %s for block:\n%s",
71                 err,
72                 formatting.DumpBytes{Bytes: blkBytes})
73             t.GetFailed(vdr, requestID, blkID)
74             return
75         }
76
77         t.insertFrom(vdr, blk)
78     }

```

Listing 3.60: snow/engine/snowman/transitive.go

Put handles an incoming Put request from other validators, it would call VM.ParseBlock to decode the Container.

```

272     // ParseBlock implements the snowman.ChainVM interface
273     func (vm *VM) ParseBlock(b []byte) (snowman.Block, error) {
274         vm.metalock.Lock()
275         defer vm.metalock.Unlock()
276
277         ethBlock := new(types.Block)
278         if err := rlp.DecodeBytes(b, ethBlock); err != nil {
279             return nil, err
280         }
281         block := &Block{
282             id:      ids.NewID(ethBlock.Hash()),
283             ethBlock: ethBlock,
284             vm:      vm,
285         }
286         vm.blockCache.Put(block.ID(), block)
287         return block, nil
288     }

```

Listing 3.61: vms/evm/vm.go

ParseBlock calls DecodeBytes to decode the raw bytes into ethBlock and then assigns it's hash to Block.id. Next, snowman checks the block's validity by calling blk.Verify().

```

270     func (t *Transitive) deliver(blk snowman.Block) {
271         if t.Consensus.Issued(blk) {
272             return
273         }
274     }

```

```

275     blkID := blk.ID()
276     t.pending.Remove(blkID)
277
278     if err := blk.Verify(); err != nil {
279         t.Config.Context.Log.Debug("Block failed verification due to %s, dropping
                block", err)
280         t.blocked.Abandon(blkID)
281         t.numBlockedBlk.Set(float64(t.pending.Len())) // Tracks performance
                statistics
282         return
283     }
284
285     t.Config.Context.Log.All("Adding block to consensus: %s", blkID)
286
287     t.Consensus.Add(blk)
288     t.pushSample(blk)

```

Listing 3.62: snow/engine/snowman/transitive.go

```

57     // Verify implements the snowman.Block interface
58     func (b *Block) Verify() error {
59         _, err := b.vm.chain.InsertChain([]*types.Block{b.ethBlock})
60         return err
61     }

```

Listing 3.63: vms/evm/block.go

Finally, InsertChain calls verifyHeaderWorker to confirm the new block's header is valid.

```

56     // modified from consensus.go
57     func (self *DummyEngine) verifyHeader(chain consensus.ChainReader, header, parent *
        types.Header, uncle bool, seal bool) error {
58         // Ensure that the header's extra-data section is of a reasonable size
59         if uint64(len(header.Extra)) > myparams.MaximumExtraDataSize {
60             return fmt.Errorf("extra-data too long: %d > %d", len(header.Extra),
                myparams.MaximumExtraDataSize)
61         }
62         // Verify the header's timestamp
63         if !uncle {
64             if header.Time > uint64(time.Now().Add(allowedFutureBlockTime).Unix()) {
65                 return consensus.ErrFutureBlock
66             }
67         }
68         //if header.Time <= parent.Time {
69         if header.Time < parent.Time {
70             return errZeroBlockTime
71         }
72         // Verify that the gas limit is <= 2^63-1
73         cap := uint64(0x7fffffffffffffff)
74         if header.GasLimit > cap {
75             return fmt.Errorf("invalid gasLimit: have %v, max %v", header.GasLimit, cap)
76         }
77         // Verify that the gasUsed is <= gasLimit

```

```

78     if header.GasUsed > header.GasLimit {
79         return fmt.Errorf("invalid gasUsed: have %d, gasLimit %d", header.GasUsed,
            header.GasLimit)
80     }
81
82     // Verify that the gas limit remains within allowed bounds
83     diff := int64(parent.GasLimit) - int64(header.GasLimit)
84     if diff < 0 {
85         diff *= -1
86     }
87     limit := parent.GasLimit / params.GasLimitBoundDivisor
88
89     if uint64(diff) >= limit || header.GasLimit < params.MinGasLimit {
90         return fmt.Errorf("invalid gas limit: have %d, want %d += %d", header.
            GasLimit, parent.GasLimit, limit)
91     }
92     // Verify that the block number is parent's +1
93     if diff := new(big.Int).Sub(header.Number, parent.Number); diff.Cmp(big.NewInt
        (1)) != 0 {
94         return consensus.ErrInvalidNumber
95     }
96     // Verify the engine specific seal securing the block
97     if seal {
98         if err := self.VerifySeal(chain, header); err != nil {
99             return err
100         }
101     }
102     return nil
103 }
104
105 func (self *DummyEngine) verifyHeaderWorker(chain consensus.ChainReader, headers []*
    types.Header, seals []bool, index int) error {
106     var parent *types.Header
107     if index == 0 {
108         parent = chain.GetHeader(headers[0].ParentHash, headers[0].Number.Uint64()
            -1)
109     } else if headers[index-1].Hash() == headers[index].ParentHash {
110         parent = headers[index-1]
111     }
112     if parent == nil {
113         return consensus.ErrUnknownAncestor
114     }
115     if chain.GetHeader(headers[index].Hash(), headers[index].Number.Uint64()) != nil
        {
116         return nil // known block
117     }
118     return self.verifyHeader(chain, headers[index], parent, false, seals[index])
119 }

```

Listing 3.64: coreth/consensus/dummy/consensus.go

However, the implementation of Ethereum Virtual Machine in Avalanche is different from the

original PoW one, it doesn't need to calculate the `Difficulty` in the header, so `VerifySeal`(line 98) can be ignored. On the other hand, since it's not based on PoW, the massive calculation of `Header.nonce` against the `MixDigest` is omitted.

```
241 func (self *DummyEngine) VerifySeal(chain consensus.ChainReader, header *types.  
    Header) error {  
242     return nil  
243 }  
244  
245 func (self *DummyEngine) Prepare(chain consensus.ChainReader, header *types.Header)  
    error {  
246     header.Difficulty = big.NewInt(1)  
247     return nil  
248 }
```

Listing 3.65: `coreth/consensus/dummy/consensus.go`

Specifically, an attacker could craft lots of valid blocks with nil txs, i.e. empty blocks and send them out.

Since these blocks are valid, they would be added into consensus, which can be used for flood attacks. On the other hand, attackers could also set `header.Time` to the maximum extent (`15 * time.Second`), so the next block would not be accepted until that time (lines 62-71), which could significantly decrease the performance.

So technically, attackers could send out lots of valid empty block within 15 seconds `header.Time` to paralyze the ETH in the AVA network.


**Recommendation** Add a sanity check on `header.Time`.



## 4 | Conclusion

In this security audit, we have analyzed the Avalanche Blockchain. During the first phase of our audit, we studied the source code and ran our in-house analyzing tools through the codebase, including areas such as P2P networking, consensus algorithm, and transaction model, etc. A list of potential issues were found, and some of them involve unusual interactions among multiple modules, therefore we developed test cases to reproduce and verify each of them. After further analysis and internal discussion, we determined that 18 issues need to be brought up and paid more attention to, which are reported in Sections [2](#) and [3](#).

Our impression through this audit is that the Avalanche Blockchain software is neatly organized and elegantly implemented and those identified issues are promptly confirmed and fixed. We'd like to commend AVA Labs for a well-done software project, and for quickly fixing issues found during the audit process. Also, as expressed in Section [1.4](#), we appreciate any constructive feedback or suggestions about this report.



## References

- [1] AVA Labs. AVA Blockchain. <https://github.com/ava-labs/>.
- [2] AVA Labs. AVA Labs. <https://avalabs.org>.
- [3] PeckShield. PeckShield Inc. <https://www.peckshield.com>.
- [4] OWASP. Risk Rating Methodology. [https://www.owasp.org/index.php/OWASP\\_Risk\\_Rating\\_Methodology](https://www.owasp.org/index.php/OWASP_Risk_Rating_Methodology).
- [5] Lcamtuf. american fuzzy lop. <http://lcamtuf.coredump.cx/afl/>.
- [6] MITRE. CWE VIEW: Development Concepts. <https://cwe.mitre.org/data/definitions/699.html>.
- [7] MITRE. CWE CATEGORY: Bad Coding Practices. <https://cwe.mitre.org/data/definitions/1006.html>.
- [8] MITRE. CWE CATEGORY: Improper Input Validation. <https://cwe.mitre.org/data/definitions/20.html>.
- [9] MITRE. CWE CATEGORY: Business Logic Errors. <https://cwe.mitre.org/data/definitions/840.html>.
- [10] MITRE. CWE-770: Allocation of Resources Without Limits or Throttling. <https://cwe.mitre.org/data/definitions/770.html>.

- [11] MITRE. CWE-1215: Input Validation Issues. <https://cwe.mitre.org/data/definitions/1215.html>.
- [12] MITRE. CWE-349: Acceptance of Extraneous Untrusted Data With Trusted Data. <https://cwe.mitre.org/data/definitions/349.html>.

